


# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering




Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>

# CS6670: Computer Vision

Noah Snavely

## Lecture 2: Image filtering



Hybrid Images, Oliva et al., <http://cvcl.mit.edu/hybridimage.htm>




# Reading

- Szeliski, Chapter 3.1-3.2

# What is an image?



# What is an image?




**We'll focus on these in this class**

**(More on this process later)**

# What is an image?

- A grid (matrix) of intensity values



=

255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255
255	255	127	145	145	175	127	127	95	47	255	255
255	255	74	127	127	127	95	95	95	47	255	255
255	255	255	74	74	74	74	74	74	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)




# What is an image?

- We can think of a (grayscale) image as a **function**,  $f$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$  (or a 2D *signal*):
  - $f(x,y)$  gives the **intensity** at position  $(x,y)$



[snoop](#)



[3D view](#)

- A **digital** image is a discrete (**sampled, quantized**) version of this function

# Image transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (linear filtering)

# Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

What's the next best thing?

# Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function




	7	

Modified image data

# Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



# Cross-correlation

Let  $F$  be the image,  $H$  be the kernel (of size  $2k+1 \times 2k+1$ ), and  $G$  be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

# Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)


$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

- Convolution is **commutative** and **associative**

# Convolution





# Mean filtering


$H$



0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F$

=

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$G$

# Linear filters: examples



Original



0	0	0
0	1	0
0	0	0



Identical image

# Linear filters: examples



Original



0	0	0
1	0	0
0	0	0



Shifted left  
By 1 pixel

# Linear filters: examples




Original



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Blur (with a mean filter)

# Linear filters: examples




Original

$$* \left( \begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \right) =$$




**Sharpening filter**  
(accentuates edges)

# Sharpening




**before**





**after**

# Smoothing with box filter revisited




# Gaussian Kernel





$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$





# Gaussian filters




$\sigma = 1$  pixel



$\sigma = 5$  pixels




$\sigma = 10$  pixels



$\sigma = 30$  pixels


# Gaussian filter

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian




- Convolving two times with Gaussian kernel of width  $\sigma$  = convolving once with kernel of width  $\sigma\sqrt{2}$

# Sharpening




**before**




**after**

# Sharpening revisited

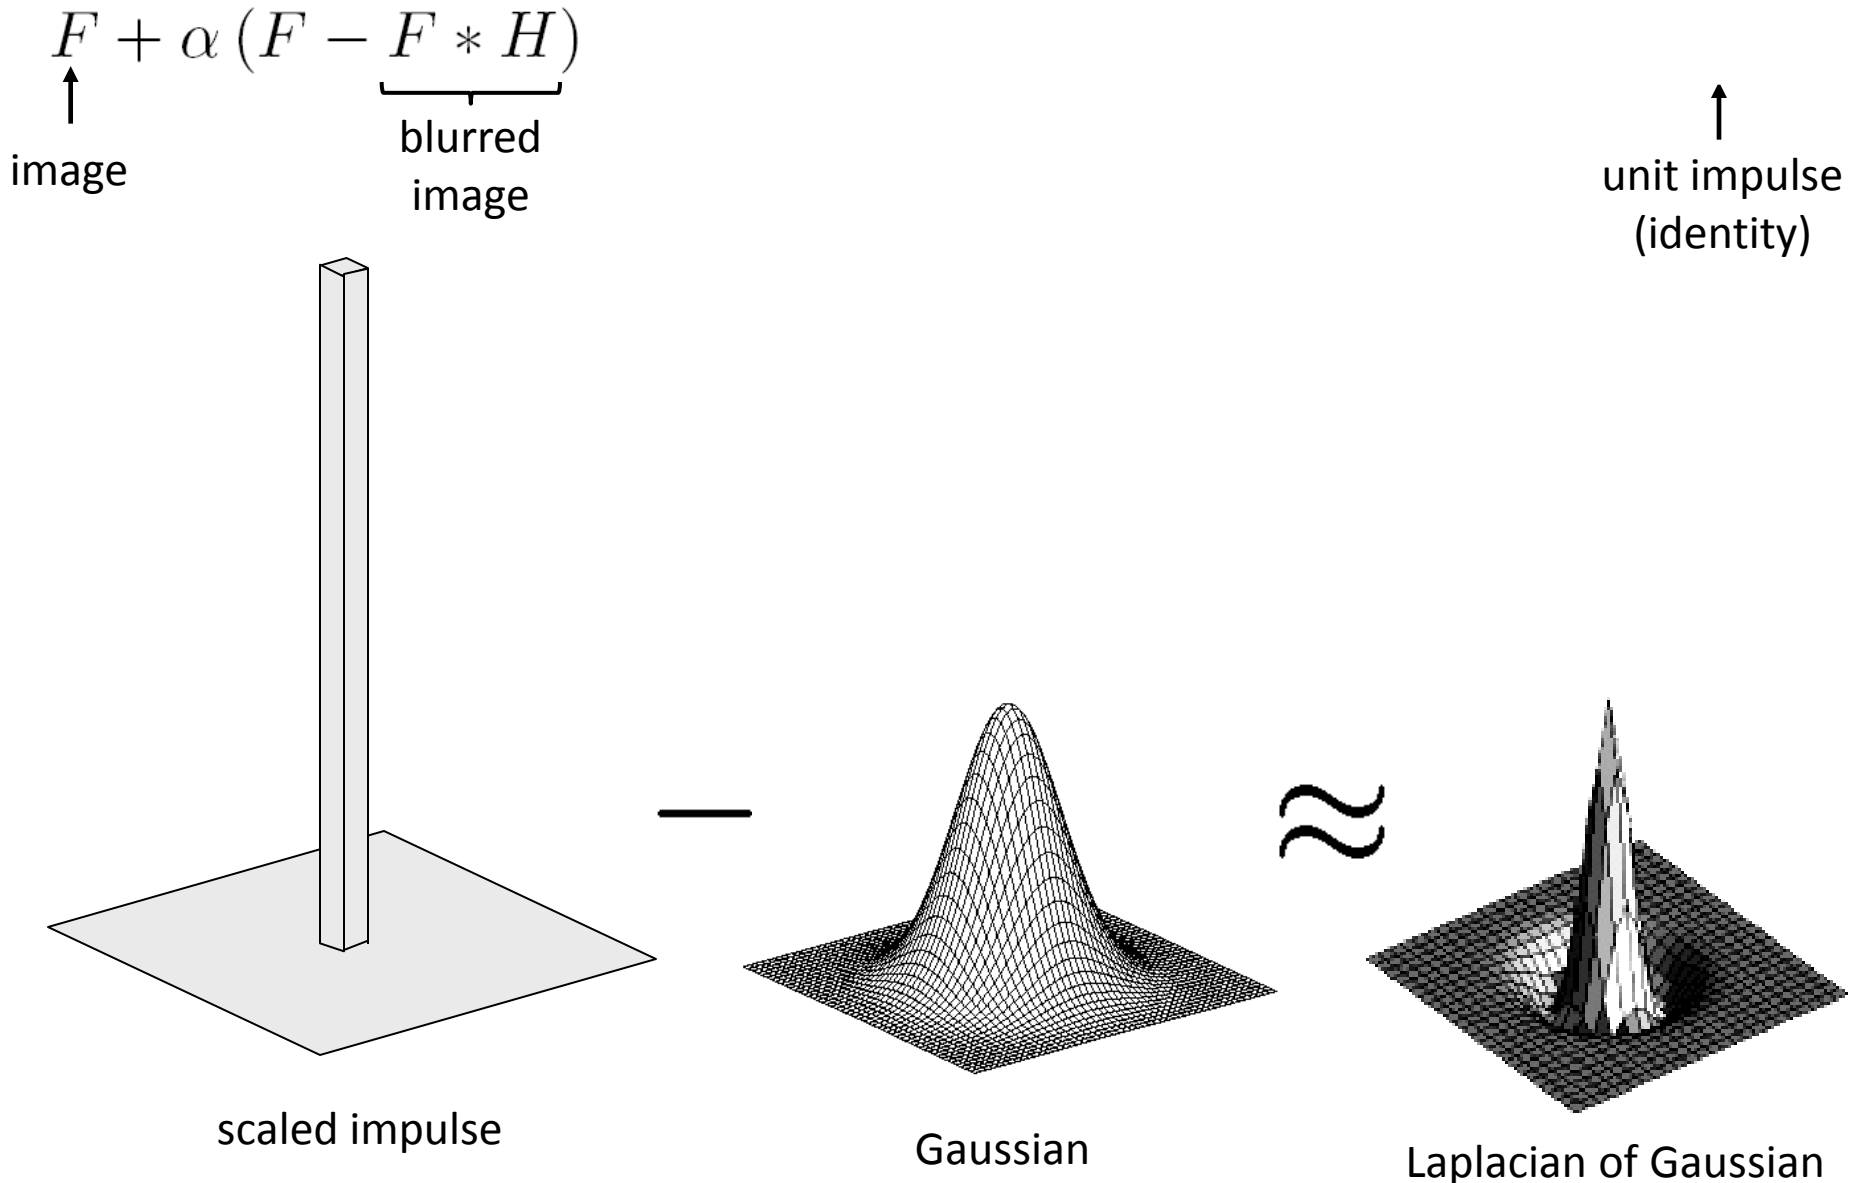
- What does blurring take away?



Let's add it back:



# Sharpen filter



# Sharpen filter





# Convolution in the real world

## Camera shake



Source: Fergus, *et al.* “Removing Camera Shake from a Single Photograph”, SIGGRAPH 2006

**Bokeh:** Blur in out-of-focus regions of an image.





Source: <http://lullaby.homepage.dk/diy-camera/bokeh.html>

# Questions?




# Edge detection



- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels


# Origin of Edges



- Edges are caused by a variety of factors

# Characterizing edges

- An edge is a place of rapid change in the image intensity function



# Image derivatives

- How can we differentiate a *digital* image  $F[x,y]$ ?
  - Option 1: reconstruct a continuous image,  $f$ , then compute the derivative
  - Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_x$




$$\frac{\partial f}{\partial y} : \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$H_y$

# Image gradient

- The *gradient* of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The *edge strength* is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:


$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?


# Image gradient




$f$



$\frac{\partial f}{\partial x}$



$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$



$\frac{\partial f}{\partial y}$

# The Sobel operator

- Common approximation of derivative of Gaussian

$$\frac{1}{8}$$

-1	0	1
-2	0	2
-1	0	1

$s_x$

$$\frac{1}{8}$$

1	2	1
0	0	0
-1	-2	-1

$s_y$

- The standard defn. of the Sobel operator omits the  $1/8$  term
  - doesn't make a difference for edge detection
  - the  $1/8$  term **is** needed to get the right gradient value

# Sobel operator: example

